




Tensorflow Basics

CS60010



Deep Learning Package Zoo



Caffe

and others ...

Deep Learning Frameworks

- Scales machine learning code
- **Computes gradients!**
- Standardise machine learning machine learning applications for sharing
- Zoo of deep learning advantages available with different advantages, levels of abstraction, programming languages, etc.
- Provides an interface with GPU for parallel processing

What is Tensorflow

 tensorflow / tensorflow

 Watch ▾

7,273

★ Star

86,389

🍴 Fork

42,106

<> Code

! Issues 1,302

🔗 Pull requests 202

📁 Projects 0

📊 Insights

Computation using data flow graphs for scalable machine learning <http://tensorflow.org>

tensorflow

machine-learning

python

deep-learning

deep-neural-networks

neural-network

ml

distributed

🕒 27,299 commits

🌿 20 branches

📦 45 releases

👤 1,245 contributors

📄 Apache-2.0


Programming Model

Idea: Express numeric computation as a *graph*

- Graph nodes are **operations** that can have any number of inputs and exactly one output
- Graph edges are **tensors** that flow between nodes

Tensors are n dimensional array


$$h = \text{ReLU}(Wx + b)$$



$$h = \text{ReLU}(Wx + b)$$

Variables are stateful nodes which output their current value.
State is retained across multiple executions of a graph

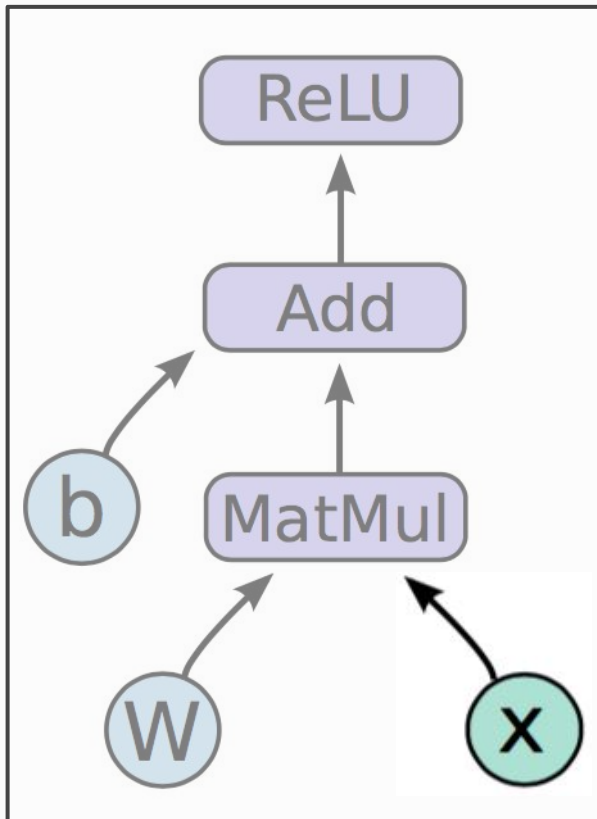
(mostly parameters)



$$h = \text{ReLU}(Wx + b)$$

Placeholders are nodes whose value is fed in at execution time

(inputs, labels, ...)




$$h = \text{ReLU}(Wx + b)$$

Mathematical operations:

MatMul: Multiply two matrix values.

Add: Add elementwise (with broadcasting).

ReLU: Activate with elementwise rectified linear function.



In code,

1. Create weights, including initialization

$W \sim \text{Uniform}(-1, 1); b = 0$

2. Create input placeholder x
 $m * 784$ input matrix

3. Build flow graph

```
import tensorflow as tf
```


```
b = tf.Variable(tf.zeros((100,)))
```

```
W = tf.Variable(tf.random_uniform((784, 100), -1, 1))
```

```
x = tf.placeholder(tf.float32, (100, 784))
```

```
h = tf.nn.relu(tf.matmul(x, W) + b)
```

$$h = \text{ReLU}(Wx + b)$$




But where is the graph?

New nodes are automatically built into the underlying graph!
`tf.get_default_graph().get_operations():`

zeros/shape
zeros/Const
zeros
Variable
Variable/Assign
Variable/read
random_uniform/shape
random_uniform/min
random_uniform/max
random_uniform/RandomUniform

random_uniform/sub
random_uniform/mul
random_uniform
Variable_1
Variable_1/Assign
Variable_1/read
Placeholder
MatMul
add
Relu == h


h refers to an op!



How do we run it?

So far we have defined a **graph**.

We can deploy this graph with a **session**:
a binding to a particular execution
context (e.g. CPU, GPU)



Getting output

```
sess.run(fetches, feeds)
```

Fetches: List of graph nodes.

Return the outputs of these nodes.

Feeds: Dictionary mapping from graph nodes to concrete values. Specifies the value of each graph node given in the dictionary.

```
import numpy as np
import tensorflow as tf
```

```
b = tf.Variable(tf.zeros((100,)))
W = tf.Variable(tf.random_uniform((784, 100),
                                  -1, 1))
```

```
x = tf.placeholder(tf.float32, (100, 784))
h = tf.nn.relu(tf.matmul(x, W) + b)
```

```
sess = tf.Session()
sess.run(tf.initialize_all_variables())
sess.run(h, {x: np.random.random(100, 784)}))
```

So what have we covered so far?

We first built a **graph** using **variables** and **placeholders**

We then deployed the graph onto a **session**, which is the **execution environment**

Next we will see how to **train** the **model**

How do we define the loss?

Use **placeholder** for **labels**

Build loss node using labels and **prediction**

```
prediction = tf.nn.softmax(...) #Output of neural network
label = tf.placeholder(tf.float32, [100, 10])

cross_entropy = -tf.reduce_sum(label * tf.log(prediction), axis=1)
```

How do we compute Gradients?

```
train_step = tf.train.GradientDescentOptimizer(0.5).minimize(cross_entropy)
```

`tf.train.GradientDescentOptimizer` is an **Optimizer** object

`tf.train.GradientDescentOptimizer(lr).minimize(cross_entropy)` adds optimization **operation** to computation graph

TensorFlow graph **nodes** have **attached gradient operations**

Gradient with respect to **parameters** computed with **backpropagation**

...automatically

Creating the train_step op

```
prediction = tf.nn.softmax(...)
label = tf.placeholder(tf.float32, [None, 10])

cross_entropy = tf.reduce_mean(-tf.reduce_sum(label * tf.log(prediction),
reduction_indices=[1]))

train_step = tf.train.GradientDescentOptimizer(0.5).minimize(cross_entropy)
```


Variable sharing: naive way

```
variables_dict = {  
    "weights": tf.Variable(tf.random_normal([784, 100]),  
                           name="weights"),  
    "biases": tf.Variable(tf.zeros([100]), name="biases")  
}
```

Not good for encapsulation!

Variable sharing: naive way

```
variables_dict = {  
    "conv1_weights": tf.Variable(tf.random_normal([5, 5, 32, 32]),  
        name="conv1_weights"),  
    "conv1_biases": tf.Variable(tf.zeros([32]), name="conv1_biases")  
    ... etc. ...  
}  
  
def my_image_filter(input_images, variables_dict):  
    conv1 = tf.nn.conv2d(input_images, variables_dict["conv1_weights"],  
        strides=[1, 1, 1, 1], padding='SAME')  
    relu1 = tf.nn.relu(conv1 + variables_dict["conv1_biases"])  
  
    conv2 = tf.nn.conv2d(relu1, variables_dict["conv2_weights"],  
        strides=[1, 1, 1, 1], padding='SAME')  
    return tf.nn.relu(conv2 + variables_dict["conv2_biases"])  
  
# The 2 calls to my_image_filter() now use the same variables  
result1 = my_image_filter(image1, variables_dict)  
result2 = my_image_filter(image2, variables_dict)
```

What's in a Name?

`tf.variable_scope()` provides simple name-spacing to avoid clashes

`tf.get_variable()` creates/accesses variables from within a variable scope

```
with tf.variable_scope("foo"):
    v = tf.get_variable("v", shape=[1]) # v.name == "foo/v:0"
```

```
with tf.variable_scope("foo", reuse=True):
    v1 = tf.get_variable("v") # Shared variable found!
```



```
with tf.variable_scope("foo", reuse=False):
    v1 = tf.get_variable("v") # CRASH foo/v:0 already exists!
```

What's in a Name?



```
def conv_relu(input, kernel_shape, bias_shape):
    # Create variable named "weights".
    weights = tf.get_variable("weights", kernel_shape,
                              initializer=tf.random_normal_initializer())
    # Create variable named "biases".
    biases = tf.get_variable("biases", bias_shape,
                              initializer=tf.constant_initializer(0.0))
    conv = tf.nn.conv2d(input, weights,
                        strides=[1, 1, 1, 1], padding='SAME')
    return tf.nn.relu(conv + biases)

def my_image_filter(input_images):
    with tf.variable_scope("conv1"):
        # Variables created here will be named "conv1/weights", "conv1/biases".
        relu1 = conv_relu(input_images, [5, 5, 32, 32], [32])
    with tf.variable_scope("conv2"):
        # Variables created here will be named "conv2/weights", "conv2/biases".
        return conv_relu(relu1, [5, 5, 32, 32], [32])
```

Noisy Accuracy Curve



Slow down



Learning rate 0.003 at start then dropping exponentially to 0.0001

Regularisation

```
t = tf.Variable(...)  
reg_loss = tf.nn.l2_loss(t, name=None)
```

Next class:

More on regularization and best practices in Tensorflow

Thanks!

References:

1. CS231n
2. CS224n
3. Martin Gorner's Slides on Tensorflow
4. [tensorflow.org](https://www.tensorflow.org)

